

## SGD Algorithm to predict movie ratings

There will be some functions that start with the word "grader" ex: grader\_matrix(), grader\_mean(), grader\_dim() etc, you should not change those function definition.

Every Grader function has to return True.

1. Download the data from [here](#)
2. The data will be of this format, each data point is represented as a triplet of user\_id, movie\_id and rating

| user_id | movie_id | rating |
|---------|----------|--------|
| 77      | 236      | 3      |
| 471     | 208      | 5      |
| 641     | 401      | 4      |
| 31      | 298      | 4      |
| 58      | 504      | 5      |
| 235     | 727      | 5      |

### Task 1

**Predict the rating for a given (user\_id, movie\_id) pair**

Predicted rating  $\hat{y}_{ij}$  for user  $i$ , movie  $j$  pair is calculated as  $\hat{y}_{ij} = \mu + b_i + c_j + u_i^T v_j$ , here we will be finding the best values of  $b_i$  and  $c_j$  using SGD algorithm with the optimization problem for  $N$  users and  $M$  movies is defined as

$$L = \min_{b, c, \{u_i\}_{i=1}^N, \{v_j\}_{j=1}^M} \alpha \left( \sum_j \sum_k v_{jk}^2 + \sum_i \sum_k u_{ik}^2 + \sum_i b_i^2 + \sum_j c_j^2 \right) + \sum_{i,j \in \mathcal{I}^{\text{train}}} (y_{ij} - \mu - b_i - c_j - u_i^T v_j)^2$$


```
In [ ]: du = alpha*U1[user_id] -2*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U[user_id], V.T[item_id]))*V[item_id]
```

- $\mu$  : scalar mean rating
- $b_i$  : scalar bias term for user  $i$
- $c_j$  : scalar bias term for movie  $j$
- $u_i$  : K-dimensional vector for user  $i$
- $v_j$  : K-dimensional vector for movie  $j$

\*. We will be giving you some functions, please write code in that functions only.

\*. After every function, we will be giving you expected output, please make sure that you get that output.

1. Construct adjacency matrix with the given data, assuming its [weighted un-directed bi-partited graph](#) and the weight of each edge is the rating given by user to the movie



you can construct this matrix like  $A[i][j] = r_{ij}$  here  $i$  is user\_id,  $j$  is movieid and  $r_{ij}$  is rating given by user  $i$  to the movie  $j$

Hint : you can create adjacency matrix using `csr_matrix`

- We will Apply SVD decomposition on the Adjacency matrix [link1](#), [link2](#) and get three matrices  $U, \Sigma, V$  such that  $U \times \Sigma \times V^T = A$ , if  $A$  is of dimensions  $N \times M$  then  
 $U$  is of  $N \times k$ ,  
 $\Sigma$  is of  $k \times k$  and  
 $V$  is  $M \times k$  dimensions.
  - \*. So the matrix  $U$  can be represented as matrix representation of users, where each row  $u_i$  represents a  $k$ -dimensional vector for a user
  - \*. So the matrix  $V$  can be represented as matrix representation of movies, where each row  $v_j$  represents a  $k$ -dimensional vector for a movie.
- Compute  $\mu$ ,  $\mu$  represents the mean of all the rating given in the dataset. (write your code in [def m\\_u\(\)](#))
- For each unique user initialize a bias value  $B_i$  to zero, so if we have  $N$  users  $B$  will be a  $N$  dimensional vector, the  $i^{th}$  value of the  $B$  will corresponds to the bias term for  $i^{th}$  user (write your code in [def initialize\(\)](#))
- For each unique movie initialize a bias value  $C_j$  zero, so if we have  $M$  movies  $C$  will be a  $M$  dimensional vector, the  $j^{th}$  value of the  $C$  will corresponds to the bias term for  $j^{th}$  movie (write your code in [def initialize\(\)](#))
- Compute  $dL/db_i$  (Write you code in [def derivative\\_db\(\)](#))
- Compute  $dL/dc_j$  (write your code in [def derivative\\_dc\(\)](#))
- Print the mean squared error with predicted ratings.

```
for each epoch:
    for each pair of (user, movie):
        b_i = b_i - learning_rate * dL/db_i
        c_j = c_j - learning_rate * dL/dc_j
    predict the ratings with formula
```

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot\_product}(u_i, v_j)$$

- you can choose any learning rate and regularization term in the range  $10^{-3}$  to  $10^2$
- bonus:** instead of using SVD decomposition you can learn the vectors  $u_i, v_j$  with the help of SGD algo similar to  $b_i$  and  $c_j$

## Task 2

As we know  $U$  is the learned matrix of user vectors, with its  $i$ -th row as the vector  $u_i$  for user  $i$ . Each row of  $U$  can be seen as a "feature vector" for a particular user.

The question we'd like to investigate is this: do our computed per-user features that are optimized for predicting movie ratings contain anything to do with gender?

The provided data file [user\\_info.csv](#) contains an `is_male` column indicating which users in the dataset are male. Can you predict this signal given the features `U`?

**Note 1** : there is no train test split in the data, the goal of this assignment is to give an intuition about how to do matrix factorization with the help of SGD and application of truncated SVD. for better understanding of the collaborative filtering please check netflix case study.

**Note 2** : Check if scaling of  $U$ ,  $V$  matrices improve the metric

### Reading the csv file

```
In [88]: import pandas as pd
data=pd.read_csv('ratings_train.csv')
data.head()
```

```
Out[88]:
```

|   | user_id | item_id | rating |
|---|---------|---------|--------|
| 0 | 772     | 36      | 3      |
| 1 | 471     | 228     | 5      |
| 2 | 641     | 401     | 4      |
| 3 | 312     | 98      | 4      |
| 4 | 58      | 504     | 5      |

### Create your adjacency matrix

```
In [89]: from scipy.sparse import csr_matrix
adjacency_matrix = csr_matrix((data.rating.values, (data.user_id.values,
data.item_id.values)),)
```

```
In [90]: adjacency_matrix.shape
```

```
Out[90]: (943, 1681)
```

### Grader function - 1

```
In [91]: def grader_matrix(matrix):
assert(matrix.shape==(943,1681))
return True
grader_matrix(adjacency_matrix)
```

```
Out[91]: True
```

### SVD decomposition

Sample code for SVD decomposition

```
In [92]: from sklearn.utils.extmath import randomized_svd
import numpy as np
matrix = np.random.random((20, 10))
U, Sigma, VT = randomized_svd(matrix, n_components=5, n_iter=5, random_state=None)
print(U.shape)
print(Sigma.shape)
print(VT.T.shape)
```

```
(20, 5)
(5,)
(10, 5)
```

```
In [93]: print(Sigma)
```

```
[6.91588035 1.73008049 1.60345043 1.4140752 1.18549129]
```

### Write your code for SVD decomposition

```
In [94]: U, Sigma, VT = randomized_svd(adjacency_matrix, n_components=5, n_iter=5, random_state=None)
print(U.shape)
print(Sigma.shape)
print(VT.T.shape)
```

```
(943, 5)
(5,)
(1681, 5)
```

### Compute mean of ratings

```
In [95]: def m_u(ratings):
          sum = 0
          for i in ratings:
              sum = sum + i

          mean = sum / len(ratings)

          '''In this function, we will compute mean for all the ratings'''
          # you can use mean() function to do this
          # check this (https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.mean.html) link for more details.

          return mean
```

```
In [96]: mu=m_u(data['rating'])
          print(mu)
```

3.529480398257623

Grader function -2

```
In [97]: def grader_mean(mu):
          assert(np.round(mu,3)==3.529)
          return True
          mu=m_u(data['rating'])
          grader_mean(mu)
```

Out[97]: True

Initialize  $B_i$  and  $C_j$ 

Hint : Number of rows of adjacent matrix corresponds to user dimensions ( $B_i$ ), number of columns of adjacent matrix corresponds to movie dimensions ( $C_j$ )

```
In [98]: def initialize(dim):
          biased = np.zeros(dim)
          '''In this function, we will initialize bias value 'B' and 'C'.'''
          # initialize the value to zeros
          # return output as a list of zeros

          return biased
```

```
In [99]: b_i=initialize(943)
```

```
In [100... c_j=initialize(1681)
```

Grader function -3

```
In [101... def grader_dim(b_i,c_j):
          assert(len(b_i)==943 and np.sum(b_i)==0)
          assert(len(c_j)==1681 and np.sum(c_j)==0)
          return True
          grader_dim(b_i,c_j)
```

Out[101... True

Compute  $dL/db_i$ 

```
In [102... def derivative_db(user_id,item_id,rating,U,V,mu,alpha):
          db = alpha*b_i[user_id] -2*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U[user_id], V.T[item_id]))
          '''In this function, we will compute dL/db_i'''
          return db
```

Grader function -4

```
In [103... def grader_db(value):
          assert(np.round(value,3)==-0.931)
          return True
          U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=24)
          # Please don't change random state
          # Here we are considering n_componets = 2 for our convinence
          alpha=0.01
          value=derivative_db(312,98,4,U1,V1,mu,alpha)
          grader_db(value)
```

Out[103... True

Compute  $dL/dc_j$

```
In [104... def derivative_dc(user_id,item_id,rating,U,V,mu, alpha):
    dc = alpha*c_j[item_id] -2*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U[user_id], V.T[item_id]))
    '''In this function, we will compute dL/dc_j'''
    return dc
```

Grader function - 5

```
In [105... def grader_dc(value):
    assert(np.round(value,3)==-2.929)
    return True
U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=24)
# Please don't change random state
# Here we are considering n_componets = 2 for our convinence
r=0.01
value=derivative_dc(58,504,5,U1,V1,mu)
grader_dc(value)
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-105-1f5ab27618cb> in <module>()
      6 # Here we are considering n_componets = 2 for our convinence
      7 r=0.01
----> 8 value=derivative_dc(58,504,5,U1,V1,mu)
      9 grader_dc(value)
```

TypeError: derivative\_dc() missing 1 required positional argument: 'alpha'

```
In [106... alpha = 0.01
value=derivative_dc(58,504,5,U1,V1,mu,alpha)
print(value)
```

-2.9290787114434913

```
In [117... #Compute gradient w.r.to 'u'
def derivative_du(user_id,item_id,rating,U,V,mu, alpha):
    du = 2*alpha*U[user_id] -2*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U[user_id], V.T[item_id]))*V.T[item_id]

    return du
```

```
In [120... #Compute gradient w.r.to 'v'
def derivative_dv(user_id,item_id,rating,U,V,mu, alpha):
    dv = 2*alpha*V.T[item_id] -2*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U[user_id], V.T[item_id]))*U[user_id]

    return dv
```

Compute MSE (mean squared error) for predicted ratings

for each epoch, print the MSE value

for each epoch:

for each pair of (user, movie):

b\_i = b\_i - learning\_rate \* dL/db\_i

c\_j = c\_j - learning\_rate \* dL/dc\_j

predict the ratings with formula

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot\_product}(u_i, v_j)$$

```
In [125... U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=10,n_iter=5, random_state=24)
alpha = 0.01
learning_rate=0.0001
epochs = 50
MSE = []
for i in range(epochs):
    for j , row in data.iterrows():
        user_id = row['user_id']
        item_id = row['item_id']
        rating = row['rating']
        db=derivative_db(user_id,item_id,rating,U1,V1,mu,alpha)
        dc=derivative_dc(user_id,item_id,rating,U1,V1,mu,alpha)

        b_i[user_id] = b_i[user_id] - learning_rate * db
        c_j[item_id] = c_j[item_id] - learning_rate * dc

        du=derivative_du(user_id,item_id,rating,U1,V1,mu,alpha)
        dv=derivative_dv(user_id,item_id,rating,U1,V1,mu,alpha)

        U1[user_id] = U1[user_id] - learning_rate * du
        V1.T[item_id] = V1.T[item_id] - learning_rate * dv
```

```

sum = 0
for j , row in data.iterrows():
    user_id = row['user_id']
    item_id = row['item_id']
    rating = row['rating']
    sum = sum + pow((rating-(mu+b_i[user_id]+c_j[item_id]+np.dot(U1[user_id], V1.T[item_id]))),2)

avg = sum/data.shape[0]
MSE.append(avg)


```

In [126...

```

epoch = np.arange(50)
import matplotlib.pyplot as plt
plt.figure(figsize=(6,5))
plt.grid()
plt.plot(epoch,MSE,color='orange')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title('Plot of Epoch Number Vs MSE',fontsize = 14)
plt.legend(['MSE loss'])
plt.show()

```



## Checking if scaling user and item vector can improve metric loss

In [195...

```

from sklearn import preprocessing
#Normalize Data
U2 = preprocessing.normalize(U1)
V2 = preprocessing.normalize(V1.T)
V2 = V2.T
V2.shape

```

Out[195...] (10, 1681)

In [196...

```

alpha = 0.01
learning_rate=0.0001
epochs = 50
MSE = []
for i in range(epochs):
    for j , row in data.iterrows():
        user_id = row['user_id']
        item_id = row['item_id']
        rating = row['rating']
        db=derivative_db(user_id,item_id,rating,U2,V2,mu,alpha)
        dc=derivative_dc(user_id,item_id,rating,U2,V2,mu,alpha)

        b_i[user_id] = b_i[user_id] - learning_rate * db
        c_j[item_id] = c_j[item_id] - learning_rate * dc

        du=derivative_du(user_id,item_id,rating,U2,V2,mu,alpha)
        dv=derivative_dv(user_id,item_id,rating,U2,V2,mu,alpha)

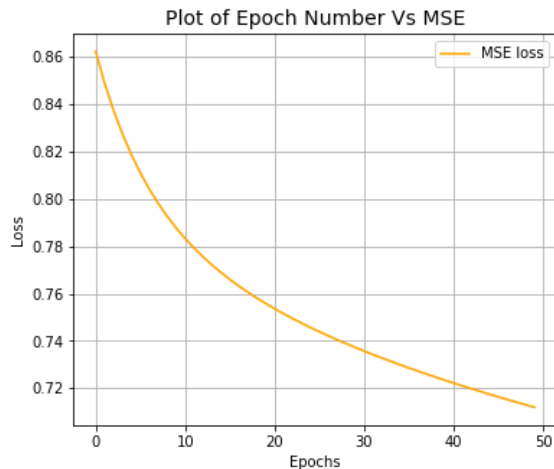
        U2[user_id] = U2[user_id] - learning_rate * du
        V2.T[item_id] = V2.T[item_id] - learning_rate * dv

    sum = 0
    for j , row in data.iterrows():
        user_id = row['user_id']
        item_id = row['item_id']
        rating = row['rating']
        sum = sum + pow((rating-(mu+b_i[user_id]+c_j[item_id]+np.dot(U2[user_id], V2.T[item_id]))),2)

```

```
avg = sum(data.shape[0])
MSE.append(avg)
```

```
In [197... epoch = np.arange(50)
import matplotlib.pyplot as plt
plt.figure(figsize=(6,5))
plt.grid()
plt.plot(epoch,MSE,color='orange')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title('Plot of Epoch Number Vs MSE',fontsize = 14)
plt.legend(['MSE loss'])
plt.show()
```



## Conclusion

metric loss is improving after scaling the user and item vector.

### Task 2

```
In [127... import pandas as pd
data1=pd.read_csv('user_info.csv.txt')
data1.head()
```

```
Out[127... user_id  age  is_male  orig_user_id
0      0    24        1          1
1      1    53        0          2
2      2    23        1          3
3      3    24        1          4
4      4    33        0          5
```

```
In [128... X = U1
Y = data1['is_male']
```

```
In [140... #Initialize weights
def initialize_weights(dim):
    w=np.zeros(dim)
    b=0
    return w,b

#Compute sigmoid
import math
def sigmoid(z):
    sig = 1/(1 + math.exp(-z))

    return sig

#Compute Loss
def logloss(y_true,y_pred):
    n = len(y_true)
    loss = 0
    for i in range(n):
        loss += (y_true[i] *math.log(y_pred[i],10)) + (1-y_true[i])*math.log((1-y_pred[i]),10)
    log_loss = -loss/n
    #'''In this function, we will compute log loss '''
```

```

    return log_loss

#Compute gradient w.r.to 'w'
def gradient_dw(x,y,w,b,alpha,N):
    y1 = np.dot(w.T,x) + b
    sig = sigmoid(y1)
    dw = x*(y - sig) -(alpha/N)*w.T

    return dw

#Compute gradient w.r.to 'b'
def gradient_db(x,y,w,b):
    y1 = np.dot(w.T,x) + b
    sig = sigmoid(y1)

    db = y - sig
    '''In this function, we will compute gradient w.r.to b '''
    return db

#Implementing Logistic regression
def train(X_train,y_train,epochs,alpha,eta0):
    train_loss = list()
    w,b = initialize_weights(X.shape[1])
    N = X.shape[0]
    for i in range(epochs):
        for i in range(len(X_train)):
            dw = gradient_dw(X_train[i],y_train[i],w,b,alpha,N)
            db = gradient_db(X_train[i],y_train[i],w,b)
            w = w + eta0*dw
            b = b + eta0*db

        y_train_pred = []
        for i in range(len(X_train)):
            z = np.dot(w.T,X_train[i]) + b
            sig = sigmoid(z)
            y_train_pred.append(sig)
        train_loss.append(logloss(y_train,y_train_pred))

    return w,b,train_loss

```

```

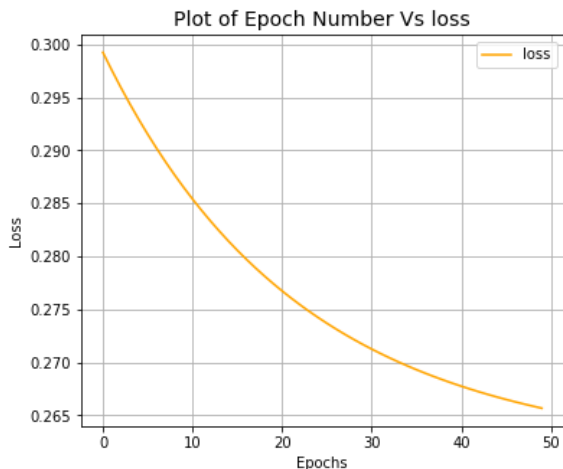
In [176... alpha=0.0001
eta0=0.0001
N=X.shape[0]
epochs=50
w,b,loss =train(X,Y,epochs,alpha,eta0)

```

```

In [177... epoch = np.arange(50)
import matplotlib.pyplot as plt
plt.figure( figsize=(6,5))
plt.grid()
plt.plot(epoch,loss,color='orange')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title('Plot of Epoch Number Vs loss',fontsize = 14)
plt.legend(['loss'])
plt.show()

```



```

In [178... def pred(w,b, X):
    N = X.shape[0]
    predict = []
    for i in range(N):

```



```
z=np.dot(w,X[i])+b
if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
    predict.append(1)
else:
    predict.append(0)
return np.array(predict)
```

```
In [179... y_pred = pred(w,b, X)
y_pred[0:50]
```

```
Out[179... array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
      1, 1, 1, 1, 1, 1])
```

## Conclusion

User feature related to ith row of user does not found to be true to predict whether the user is Male or Female.